

Homework 6_1

Jyoti Chaudhary

March 22, 2018

Solution 1.a

```
# Fitting the MLR model using LM function
```

```
data_B15 <- read.csv(paste(getwd(), "/data_table_B15.csv", sep = ""), header=T)
```

```
colnames(data_B15) <- c('x6', 'y', 'x1', 'x2', 'x3', 'x4', 'x5')
```

```
data_B15_fit <- lm(formula = data_B15$y ~ data_B15$x1 + data_B15$x2 + data_B15$x3 + data_B15$x4  
+ data_B15$x5, data = data_B15)
```

```
summary(data_B15_fit)
```

```
##  
## Call:  
## lm(formula = data_B15$y ~ data_B15$x1 + data_B15$x2 + data_B15$x3 +  
##     data_B15$x4 + data_B15$x5, data = data_B15)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -91.38 -18.97  -3.56   16.00   91.83   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  995.63646    91.64099   10.865 3.35e-15 ***  
## data_B15$x1    1.40734     0.68914    2.042 0.046032 *    
## data_B15$x2  -14.80139     7.02747   -2.106 0.039849 *    
## data_B15$x3    3.19909     0.62231    5.141 3.89e-06 ***  
## data_B15$x4   -0.10797     0.13502   -0.800 0.427426      
## data_B15$x5    0.35518     0.09096    3.905 0.000264 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 37.09 on 54 degrees of freedom  
## Multiple R-squared:  0.6746, Adjusted R-squared:  0.6444   
## F-statistic: 22.39 on 5 and 54 DF,  p-value: 4.407e-12
```

The MLR equation from above R output is:

$$\hat{y} = 996 + 1.41 * x_1 - 14.8 * x_2 + 3.2 * x_3 - 0.10 * x_4 + 0.3552 * x_5$$

where

City = x6 Mort = y PRECIP = x1 EDUC = x2 Nonwhite = x3 NOX = x4 SO2 = x5

Solution (1.b)

$H_0: \beta_1 = \beta_2 = \beta_3 = \beta_4 = \beta_5 = 0$ vs. H_1 : at least one of 5 is non-zero.

From above, it can be seen that F-statistic: 22.39 with p-value approximately equals 0.000. So, we reject the null hypothesis

So, the data provides sufficient evidence to support the fact that the regression model is significant.

Solution (1.c)

Use t tests to assess the contribution of each regressor to the model. Discuss your findings.

From above model fit summary statistics, it can be seen that $\beta_4 = \text{NOX}$ has a high p-value of approx = 0.43. β_4 is barely significant at 5% level but not for any $\alpha < 0.427$. All other predictors have very low p-value, hence all other predictors are significant for smaller values of α .

PRECIP (x1), EDUC(x2), NONWHITE(x3) and SO2(x5) contribute to the model.

solution (1.d)

Find a 95% CI for the regression coefficient for SO2.

```
confint(data_B15_fit, "data_B15$x5", level = 0.95)
```

```
##                2.5 %    97.5 %  
## data_B15$x5 0.1728118 0.5375405
```

From above output, the 95% C.I for SO2 is (0.1728, 0.5375)

solution (1.e)

Run all possible models and choose the best one with justifications.

```
fit_all =regsubsets(y ~ x1 + x2 + x3 + x4 + x5, data=data_B15, nbest=10,really.big=T, intercept  
=T)
```

```
model_fit_all =summary(fit_all)[[1]]  
RSQ =summary(fit_all)[[2]]  
SSE=summary(fit_all)[[3]]  
ADJR2 =summary(fit_all)[[4]]  
Cp=summary(fit_all)[[5]]  
BIC=summary(fit_all)[[6]]
```

```
fit_statistics =cbind(model_fit_all,RSQ,SSE,ADJR2,Cp,BIC)  
fit_statistics
```

##	(Intercept)	x1	x2	x3	x4	x5	RSQ	SSE	ADJR2	Cp
## 1	1	0	0	1	0	0	0.419250810	132570.75	0.40923789	40.364406
## 1	1	0	1	0	0	0	0.261104467	168671.67	0.24836489	66.605818
## 1	1	1	0	0	0	0	0.259582453	169019.11	0.24681663	66.858367
## 1	1	0	0	0	0	1	0.181435802	186858.07	0.16732263	79.825334
## 1	1	0	0	0	1	0	0.005731729	226966.98	-0.01141083	108.980121
## 2	1	0	1	1	0	0	0.566777866	98893.95	0.55157709	17.885066
## 2	1	0	0	1	0	1	0.525110918	108405.49	0.50844814	24.798912
## 2	1	1	0	0	0	1	0.493064509	115720.90	0.47527730	30.116411
## 2	1	1	0	1	0	0	0.492959663	115744.83	0.47516877	30.133808
## 2	1	0	0	1	1	0	0.427758360	130628.69	0.40767971	40.952738
## 2	1	0	1	0	0	1	0.360314920	146024.36	0.33786983	52.143709
## 2	1	1	1	0	0	0	0.349356056	148526.00	0.32652644	53.962126
## 2	1	1	0	0	1	0	0.298159792	160212.85	0.27353382	62.457183
## 2	1	0	1	0	1	0	0.262948911	168250.63	0.23708747	68.299767
## 2	1	0	0	0	1	1	0.257386457	169520.40	0.23132984	69.222752
## 3	1	1	0	1	0	1	0.640564447	82050.29	0.62130897	7.641570
## 3	1	0	1	1	0	1	0.629800512	84507.43	0.60996840	9.427643
## 3	1	0	0	1	1	1	0.587097292	94255.53	0.56497750	16.513439
## 3	1	1	1	1	0	0	0.578671932	96178.83	0.55610079	17.911469
## 3	1	0	1	1	1	0	0.566777911	98893.94	0.54356959	19.885058
## 3	1	1	1	0	0	1	0.514980267	110718.07	0.48899707	28.479903
## 3	1	1	0	1	1	0	0.497230898	114769.82	0.47029684	31.425077
## 3	1	1	0	0	1	1	0.493080859	115717.17	0.46592448	32.113698
## 3	1	1	1	0	1	0	0.386682523	140005.29	0.35382623	49.768501
## 3	1	0	1	0	1	1	0.376902004	142237.94	0.34352175	51.391394
## 4	1	1	1	1	0	1	0.670710400	75168.71	0.64676207	4.639416
## 4	1	0	1	1	1	1	0.649430493	80026.39	0.62393453	8.170417
## 4	1	1	0	1	1	1	0.647828947	80391.99	0.62221651	8.436163
## 4	1	1	1	1	1	0	0.582676441	95264.70	0.55232564	19.246996
## 4	1	1	1	0	1	1	0.515301392	110644.77	0.48005058	30.426618
## 5	1	1	1	1	1	1	0.674563903	74289.05	0.64443093	6.000000

BIC

## 1	-24.417489
## 1	-9.967235
## 1	-9.843771
## 1	-3.823518
## 1	7.843796
## 2	-37.907247
## 2	-32.397407
## 2	-28.479258
## 2	-28.466849
## 2	-21.208602
## 2	-14.523724
## 2	-13.504530
## 2	-8.959938
## 2	-6.022850
## 2	-5.571736
## 3	-45.015845
## 3	-43.245418
## 3	-36.695219
## 3	-35.483231
## 3	-33.812908

```
## 3 -27.036564
## 3 -24.880077
## 3 -24.386848
## 3 -12.954976
## 3 -12.005710
## 4 -46.177338
## 4 -42.420054
## 4 -42.146574
## 4 -31.961883
## 4 -22.981958
## 5 -42.789282
```

We can see from above result that the model that excludes NOX is the best model:

MORT (y) ~ PRECIP (x1) + EDUC(x2) + NONWHITE(x3) + SO2(x5)

This is based on this model's highest Cp value, highest Adjusted R-square and lowest BIC value.

As we had determined in solution (1.c), we predictor x4(NOX) is not a significant parameter, hence removing this predictor resulted in a better model fit.

Solution (1.f)

Run forward, backward and stepwise regression on the data.

```
fit1 <- lm(y ~ x1 + x2 + x3 + x4 + x5, data=data_B15)

# forward selection
step(lm(y ~ 1, data=data_B15), scope=list(lower=~1, upper=fit1), direction='forward')
```

```
## Start: AIC=496.64
## y ~ 1
##
##      Df Sum of Sq  RSS   AIC
## + x3    1    95705 132571 466.03
## + x2    1    59604 168672 480.48
## + x1    1    59256 169019 480.61
## + x5    1    41417 186858 486.63
## <none>          228275 496.64
## + x4    1     1308 226967 498.29
##
## Step: AIC=466.03
## y ~ x3
##
##      Df Sum of Sq  RSS   AIC
## + x2    1     33677  98894 450.45
## + x5    1     24165 108405 455.96
## + x1    1     16826 115745 459.89
## <none>          132571 466.03
## + x4    1      1942 130629 467.15
##
## Step: AIC=450.45
## y ~ x3 + x2
##
##      Df Sum of Sq  RSS   AIC
## + x5    1   14386.5  84507 443.02
## <none>          98894 450.45
## + x1    1     2715.1  96179 450.78
## + x4    1         0.0  98894 452.45
##
## Step: AIC=443.02
## y ~ x3 + x2 + x5
##
##      Df Sum of Sq  RSS   AIC
## + x1    1    9338.7  75169 437.99
## + x4    1    4481.0  80026 441.75
## <none>          84507 443.02
##
## Step: AIC=437.99
## y ~ x3 + x2 + x5 + x1
##
##      Df Sum of Sq  RSS   AIC
## <none>          75169 437.99
## + x4    1     879.66  74289 439.28
```

```
##
## Call:
## lm(formula = y ~ x3 + x2 + x5 + x1, data = data_B15)
##
## Coefficients:
## (Intercept)          x3          x2          x5          x1
##   995.8224      3.0998    -15.5697      0.3263      1.6350
```

The above output shows the best model as per Forward selection. The best model with lowest AIC value (437.99) is:

$y \sim x_3 + x_2 + x_5 + x_1$

```
# Backward Selection
step(lm(y ~ x1 + x2 + x3 + x4 + x5, data=data_B15), scope=list(lower=~1, upper=fit1), direction=
'backward', trace=FALSE)
```

```
##
## Call:
## lm(formula = y ~ x1 + x2 + x3 + x5, data = data_B15)
##
## Coefficients:
## (Intercept)          x1          x2          x3          x5
##    995.8224     1.6350    -15.5697     3.0998     0.3263
```

The above output shows the best model as per Backward selection. (Trace=FALSE removes the other models from the output and only display the best model)

```
# stepwise selection
```

```
step(lm(y ~ x1 + x2 + x3 + x4 + x5, data=data_B15), scope=list(lower=~1, upper=fit1), direction=
'both', trace=FALSE)
```

```
##
## Call:
## lm(formula = y ~ x1 + x2 + x3 + x5, data = data_B15)
##
## Coefficients:
## (Intercept)          x1          x2          x3          x5
##    995.8224     1.6350    -15.5697     3.0998     0.3263
```

Same model is selected as the best model from Forward, Backward and Stepwise selection.

```
# Mallows Cp criteria for model selection
```

```
leaps(x=data_B15[,3:7], y=data_B15[,2], names=names(data_B15[,3:7]), method="Cp")
```

```

## $which
##      x1      x2      x3      x4      x5
## 1 FALSE FALSE  TRUE FALSE FALSE
## 1 FALSE  TRUE FALSE FALSE FALSE
## 1  TRUE FALSE FALSE FALSE FALSE
## 1 FALSE FALSE FALSE FALSE  TRUE
## 1 FALSE FALSE FALSE  TRUE FALSE
## 2 FALSE  TRUE  TRUE FALSE FALSE
## 2 FALSE FALSE  TRUE FALSE  TRUE
## 2  TRUE FALSE FALSE FALSE  TRUE
## 2  TRUE FALSE  TRUE FALSE FALSE
## 2 FALSE FALSE  TRUE  TRUE FALSE
## 2 FALSE  TRUE FALSE FALSE  TRUE
## 2  TRUE  TRUE FALSE FALSE FALSE
## 2  TRUE FALSE FALSE  TRUE FALSE
## 2 FALSE  TRUE FALSE  TRUE FALSE
## 2 FALSE FALSE FALSE  TRUE  TRUE
## 3  TRUE FALSE  TRUE FALSE  TRUE
## 3 FALSE  TRUE  TRUE FALSE  TRUE
## 3 FALSE FALSE  TRUE  TRUE  TRUE
## 3  TRUE  TRUE  TRUE FALSE FALSE
## 3 FALSE  TRUE  TRUE  TRUE FALSE
## 3  TRUE  TRUE FALSE FALSE  TRUE
## 3  TRUE FALSE  TRUE  TRUE FALSE
## 3  TRUE FALSE FALSE  TRUE  TRUE
## 3  TRUE  TRUE FALSE  TRUE FALSE
## 3 FALSE  TRUE FALSE  TRUE  TRUE
## 4  TRUE  TRUE  TRUE FALSE  TRUE
## 4 FALSE  TRUE  TRUE  TRUE  TRUE
## 4  TRUE FALSE  TRUE  TRUE  TRUE
## 4  TRUE  TRUE  TRUE  TRUE FALSE
## 4  TRUE  TRUE FALSE  TRUE  TRUE
## 5  TRUE  TRUE  TRUE  TRUE  TRUE
##
## $label
## [1] "(Intercept)" "x1"          "x2"          "x3"          "x4"
## [6] "x5"
##
## $size
## [1] 2 2 2 2 2 3 3 3 3 3 3 3 3 3 3 4 4 4 4 4 4 4 4 4 4 5 5 5 5 5 6
##
## $Cp
## [1] 40.364406 66.605818 66.858367 79.825334 108.980121 17.885066
## [7] 24.798912 30.116411 30.133808 40.952738 52.143709 53.962126
## [13] 62.457183 68.299767 69.222752 7.641570 9.427643 16.513439
## [19] 17.911469 19.885058 28.479903 31.425077 32.113698 49.768501
## [25] 51.391394 4.639416 8.170417 8.436163 19.246996 30.426618
## [31] 6.000000

```

As per above output, the best model is “4 TRUE TRUE TRUE FALSE TRUE” which is the following model: $y \sim x1 + x2 + x3 + x5$ This is the same model as chosen by all the above selection methods.

```
# Ajusted R-square criteria for model selection
```

```
leaps(x=data_B15[,3:7], y=data_B15[,2], names=names(data_B15[,3:7]), method="adjr2")
```

```
## $which
```

```
##      x1      x2      x3      x4      x5
```

```
## 1 FALSE FALSE  TRUE FALSE FALSE
```

```
## 1 FALSE  TRUE FALSE FALSE FALSE
```

```
## 1  TRUE FALSE FALSE FALSE FALSE
```

```
## 1 FALSE FALSE FALSE FALSE  TRUE
```

```
## 1 FALSE FALSE FALSE  TRUE FALSE
```

```
## 2 FALSE  TRUE  TRUE FALSE FALSE
```

```
## 2 FALSE FALSE  TRUE FALSE  TRUE
```

```
## 2  TRUE FALSE FALSE FALSE  TRUE
```

```
## 2  TRUE FALSE  TRUE FALSE FALSE
```

```
## 2 FALSE FALSE  TRUE  TRUE FALSE
```

```
## 2 FALSE  TRUE FALSE FALSE  TRUE
```

```
## 2  TRUE  TRUE FALSE FALSE FALSE
```

```
## 2  TRUE FALSE FALSE  TRUE FALSE
```

```
## 2 FALSE  TRUE FALSE  TRUE FALSE
```

```
## 2 FALSE FALSE FALSE  TRUE  TRUE
```

```
## 3  TRUE FALSE  TRUE FALSE  TRUE
```

```
## 3 FALSE  TRUE  TRUE FALSE  TRUE
```

```
## 3 FALSE FALSE  TRUE  TRUE  TRUE
```

```
## 3  TRUE  TRUE  TRUE FALSE FALSE
```

```
## 3 FALSE  TRUE  TRUE  TRUE FALSE
```

```
## 3  TRUE  TRUE FALSE FALSE  TRUE
```

```
## 3  TRUE FALSE  TRUE  TRUE FALSE
```

```
## 3  TRUE FALSE FALSE  TRUE  TRUE
```

```
## 3  TRUE  TRUE FALSE  TRUE FALSE
```

```
## 3 FALSE  TRUE FALSE  TRUE  TRUE
```

```
## 4  TRUE  TRUE  TRUE FALSE  TRUE
```

```
## 4 FALSE  TRUE  TRUE  TRUE  TRUE
```

```
## 4  TRUE FALSE  TRUE  TRUE  TRUE
```

```
## 4  TRUE  TRUE  TRUE  TRUE FALSE
```

```
## 4  TRUE  TRUE FALSE  TRUE  TRUE
```

```
## 5  TRUE  TRUE  TRUE  TRUE  TRUE
```

```
##
```

```
## $label
```

```
## [1] "(Intercept)" "x1"          "x2"          "x3"          "x4"
```

```
## [6] "x5"
```

```
##
```

```
## $size
```

```
## [1] 2 2 2 2 2 3 3 3 3 3 3 3 3 3 4 4 4 4 4 4 4 4 5 5 5 5 5 6
```

```
##
```

```
## $adjr2
```

```
## [1] 0.40923789 0.24836489 0.24681663 0.16732263 -0.01141083
```

```
## [6] 0.55157709 0.50844814 0.47527730 0.47516877 0.40767971
```

```
## [11] 0.33786983 0.32652644 0.27353382 0.23708747 0.23132984
```

```
## [16] 0.62130897 0.60996840 0.56497750 0.55610079 0.54356959
```

```
## [21] 0.48899707 0.47029684 0.46592448 0.35382623 0.34352175
```

```
## [26] 0.64676207 0.62393453 0.62221651 0.55232564 0.48005058
```

```
## [31] 0.64443093
```


The model with highest Adjusted R-square value is the best model. Again the highest R-square value (0.64676207) is for model "4 TRUE TRUE TRUE FALSE TRUE" which is actually ($y \sim x_1 + x_2 + x_3 + x_5$) model.

Solution (1.g)

Do all 3 procedures picked the same model? If yes: Should it happen all the time, If NO: Why don't they pick the same?

Yes, all 3 procedures picked the same model - $MORT(y) \sim PRECIP(x_1) + EDUC(x_2) + NONWHITE(x_3) + SO_2(x_5)$

But it does not happen all the time. These 3 methods can pick different models based on various criteria. In this particular case, NOX is a weak parameter and its P-value is lesser at all steps in all three methods. In case of another dataset, the P-value of a variable may vary at different steps and can result in different model selection.

solution (1.h)

Perform the residual analysis of your final model and provide the final estimated model

```
fit2 <- lm(y ~ x1 + x2 + x3 + x5, data=data_B15)
```

```
#e. Residual analysis
```

```
e=residuals(fit2)          ## RESIDUAL
```

```
std_e=stdres(fit2)  ## STANDARDIZED RESIDUAL
```

```
std_e
```

```
##          1          2          3          4          5
## -1.3598827028 -1.4759548626 -0.2631211982 -2.6804192907 -0.2362788337
##          6          7          8          9         10
## -1.5354005019 -2.6872066609 -0.8883417599 -0.6790199876  0.5066891711
##         11         12         13         14         15
## -0.6168068923 -0.6995926630 -0.3668076236 -0.0283660424 -0.1857097341
##         16         17         18         19         20
##  0.3345870951 -0.6273781409 -0.2651710415  0.1679042739 -1.2715134876
##         21         22         23         24         25
##  0.0001167551 -0.0834399193  0.3649333526 -0.0507272184 -0.0488352698
##         26         27         28         29         30
## -0.0464615290  0.7504235322  0.4216126707  0.5217288726  0.5653351100
##         31         32         33         34         35
## -0.1316212571  1.3881670282 -0.1798955073 -0.6364588804  0.3850702058
##         36         37         38         39         40
##  1.1704760104 -0.2379237115 -1.1547469007  0.9661580479 -0.5039459305
##         41         42         43         44         45
##  0.7528619474 -0.5548814528 -0.0968002716  1.5201696880 -0.2650636596
##         46         47         48         49         50
##  1.0737264503 -0.1464990085 -0.2868016907  0.7307738731  2.5441693843
##         51         52         53         54         55
##  2.0214776153  1.6258751359 -1.0042745403  0.2028608029  0.2439231350
##         56         57         58         59         60
##  0.1109301454  0.6967063749 -1.0739814690  0.3446762675  2.4206296501
```

From above, none of the standardized residuals are greater than 3 therefore none of them are outliers.

```
r=studres(fit2)          ## STUDENTIZED RESIDUAL
max(abs(r))
```

```
## [1] 2.8568
```

The studentized residuals also don't show any outlier. However the 50th and 60th point have values greater than 2.5 and hence these points need closer analysis.

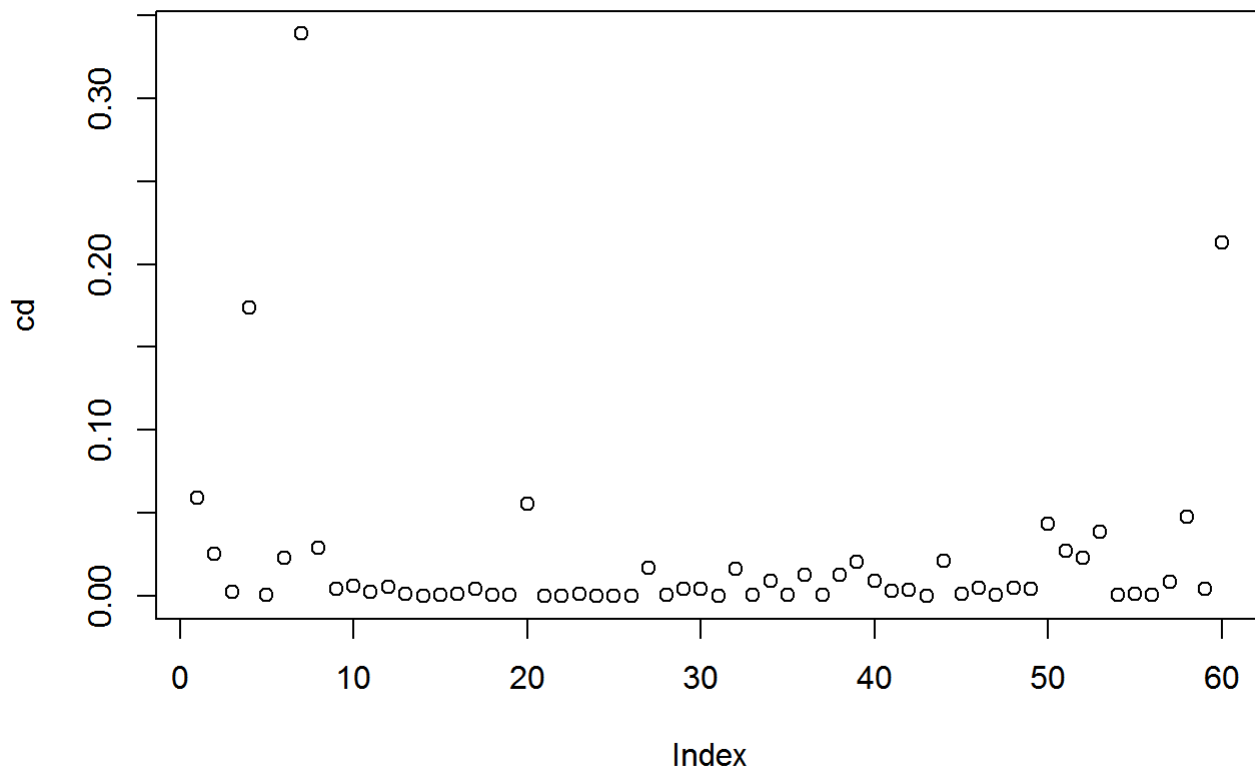
```
# Influence analysis
```

```
## COOKS DISTANCE
cd=cooks.distance(fit2)
cd
```

```
##          1          2          3          4          5
## 5.920384e-02 2.566065e-02 2.748780e-03 1.735884e-01 7.104871e-04
##          6          7          8          9         10
## 2.293407e-02 3.388075e-01 2.896764e-02 4.342522e-03 6.385758e-03
##          11         12         13         14         15
## 2.665218e-03 5.301815e-03 1.155392e-03 1.213187e-05 5.964713e-04
##          16         17         18         19         20
## 1.537115e-03 4.445000e-03 9.679526e-04 6.784888e-04 5.567658e-02
##          21         22         23         24         25
## 1.837655e-10 4.696907e-05 1.381437e-03 1.303433e-05 2.715753e-05
##          26         27         28         29         30
## 2.400820e-05 1.678024e-02 9.548235e-04 4.317069e-03 4.089604e-03
##          31         32         33         34         35
## 3.716435e-04 1.621505e-02 4.469483e-04 8.889603e-03 6.311853e-04
##          36         37         38         39         40
## 1.275908e-02 5.785972e-04 1.299847e-02 2.032373e-02 9.331727e-03
##          41         42         43         44         45
## 3.273842e-03 3.942597e-03 1.462202e-04 2.142370e-02 1.242379e-03
##          46         47         48         49         50
## 4.992725e-03 6.346837e-04 4.970891e-03 4.378052e-03 4.329587e-02
##          51         52         53         54         55
## 2.729008e-02 2.331134e-02 3.855466e-02 6.497139e-04 1.160931e-03
##          56         57         58         59         60
## 7.844472e-04 8.622769e-03 4.793206e-02 4.095782e-03 2.132721e-01
```

```
plot(cd,main="plot of cook's distance")
```

plot of cook's distance



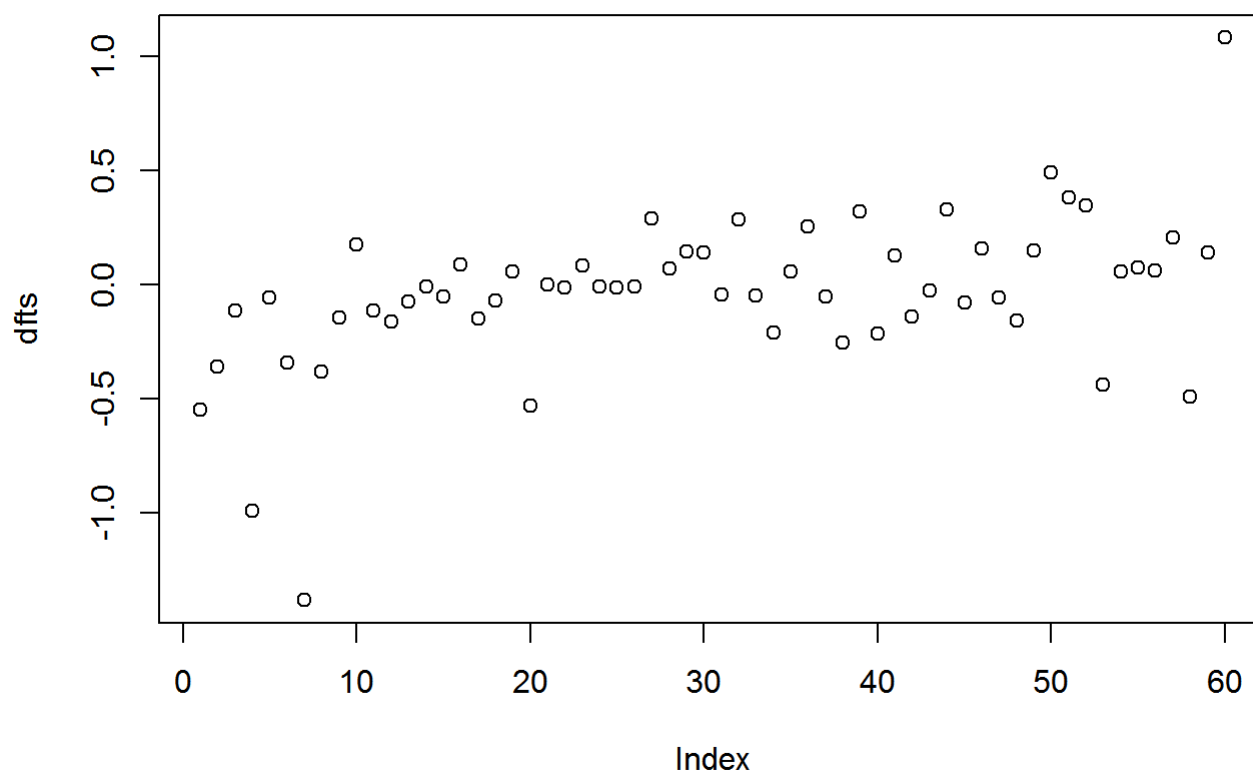
There is no point which has Cook's distance greater than 2. Therefore, as per cook's distance, no point unduly impacts the regression coefficients.

```
## DFFITS  
dfts=dffits(fit2)  
dfts
```

##	1	2	3	4	5
##	-5.484061e-01	-3.621680e-01	-1.162369e-01	-9.900532e-01	-5.908798e-02
##	6	7	8	9	10
##	-3.429684e-01	-1.383695e+00	-3.798353e-01	-1.466222e-01	1.774691e-01
##	11	12	13	14	15
##	-1.147821e-01	-1.620518e-01	-7.540447e-02	-7.717339e-03	-5.412919e-02
##	16	17	18	19	20
##	8.695531e-02	-1.482504e-01	-6.897717e-02	5.772756e-02	-5.306591e-01
##	21	22	23	24	25
##	3.003534e-05	-1.518567e-02	8.245036e-02	-7.999355e-03	-1.154663e-02
##	26	27	28	29	30
##	-1.085647e-02	2.884924e-01	6.857488e-02	1.459393e-01	1.421041e-01
##	31	32	33	34	35
##	-4.272009e-02	2.872131e-01	-4.685511e-02	-2.096751e-01	5.573977e-02
##	36	37	38	39	40
##	2.534473e-01	-5.332272e-02	-2.557268e-01	3.185806e-01	-2.145293e-01
##	41	42	43	44	45
##	1.274321e-01	-1.395117e-01	-2.679423e-02	3.313361e-01	-7.814571e-02
##	46	47	48	49	50
##	1.582230e-01	-5.582953e-02	-1.563301e-01	1.473193e-01	4.908090e-01
##	51	52	53	54	55
##	3.804235e-01	3.467209e-01	-4.390944e-01	5.649684e-02	7.553336e-02
##	56	57	58	59	60
##	6.206274e-02	2.066564e-01	-4.902481e-01	1.419510e-01	1.082503e+00

```
plot(dfts,main="plot of dffits")
```

plot of dffits



```
#threshold value  
#  $2 \times \text{square-root}(p/n) = 0.52$  ( $p = 4$ ,  $n = 60$ )
```

The 50th and 60th point have a dffits value greater than the cutoff value of 0.52. Therefore according to this analysis, these 2 data points unduly impact the parameter estimation.

```
# removing data points - 50th and 60th and see if there is improvement in model fit  
  
data1 <- data_B15[-c(1,6)]  
  
data2 <- data1[-50,]  
data3 <- data2[-59,]  
  
fit_subset <- lm(y ~ ., data = data3)  
  
summary(fit2)
```

```
##
## Call:
## lm(formula = y ~ x1 + x2 + x3 + x5, data = data_B15)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -93.600 -20.499  -2.443  17.891  92.521
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  995.82238    91.33980   10.902 2.31e-15 ***
## x1             1.63505     0.62550    2.614 0.011522 *
## x2            -15.56968     6.93862   -2.244 0.028883 *
## x3             3.09979     0.60779    5.100 4.33e-06 ***
## x5             0.32634     0.08323    3.921 0.000247 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 36.97 on 55 degrees of freedom
## Multiple R-squared:  0.6707, Adjusted R-squared:  0.6468
## F-statistic: 28.01 on 4 and 55 DF,  p-value: 1.052e-12
```

```
summary(fit_subset)
```

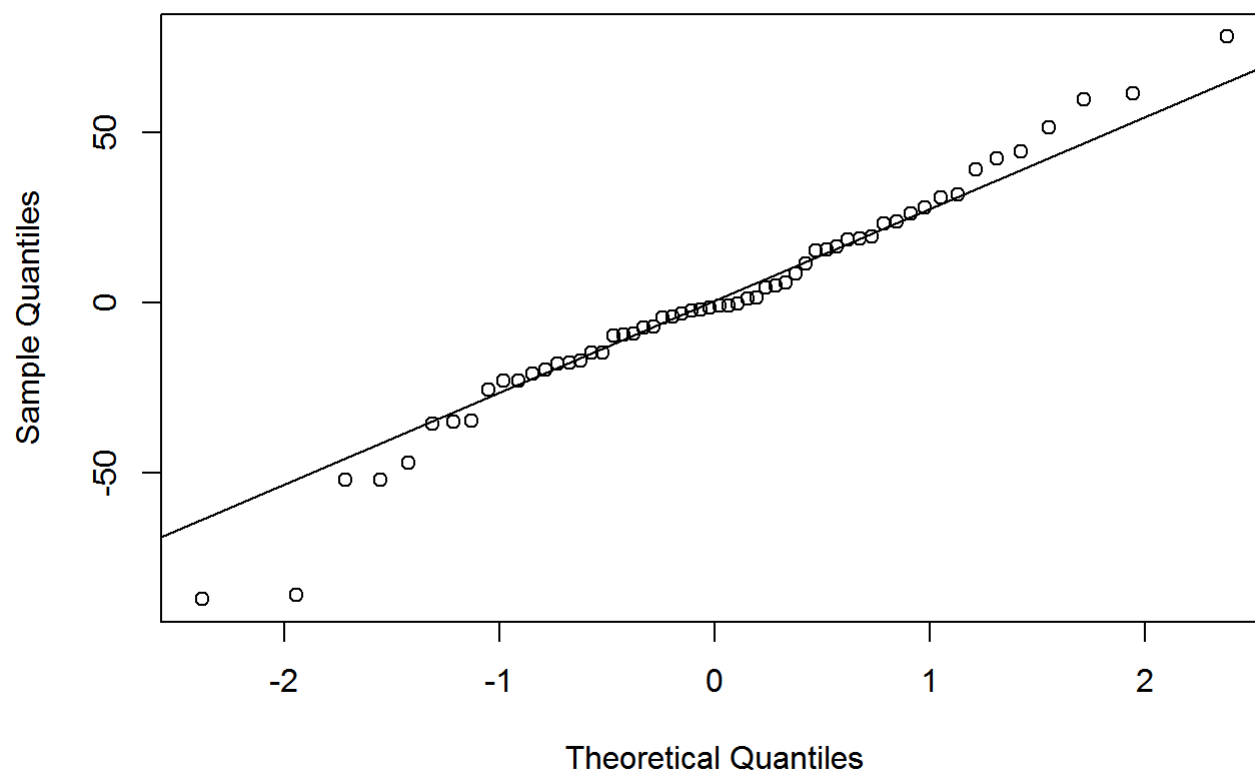
```
##
## Call:
## lm(formula = y ~ ., data = data3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -86.926 -17.467  -1.177  18.878  78.041
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  958.70795    82.93197   11.560 4.16e-16 ***
## x1             1.63894     0.56235    2.914  0.00521 **
## x2            -12.45716     6.32690   -1.969  0.05420 .
## x3             2.87892     0.57029    5.048 5.60e-06 ***
## x5             0.36838     0.07628    4.829 1.21e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 33.23 on 53 degrees of freedom
## Multiple R-squared:  0.6986, Adjusted R-squared:  0.6759
## F-statistic: 30.71 on 4 and 53 DF,  p-value: 3.074e-13
```

Removal of 50th and 60th data point has resulted in decrease in Residual standard error and increase in Adjusted R-square. Hence these data points should be excluded to attain better model fitting.

```
# Normality plot of residuals
```

```
qqnorm(fit_subset$residuals)  
qqline(fit_subset$residuals)
```

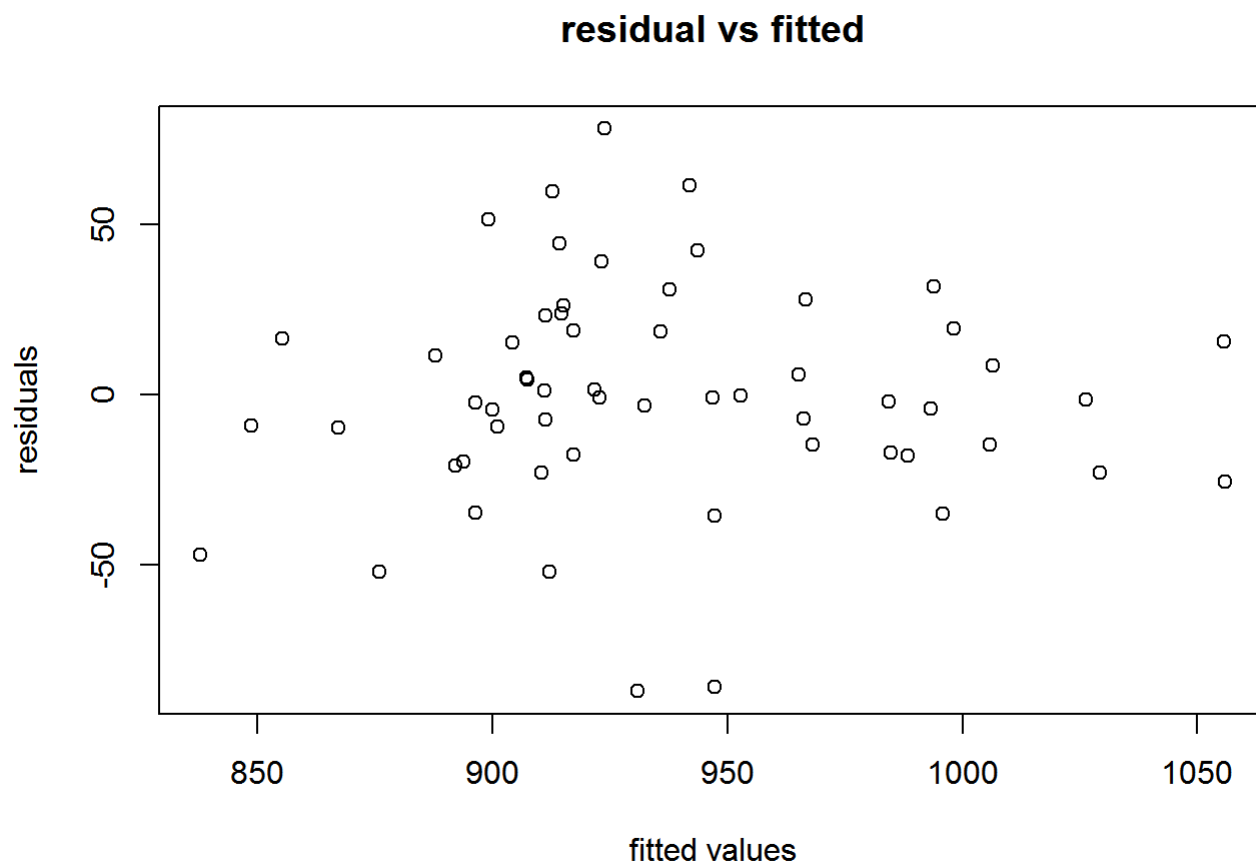
Normal Q-Q Plot



The normal probability plot shows that the residuals are much more closer to a normal distribution and there is no departure from the straight line towards the tails. This indicates that errors are IID normal.

```
# residual vs fitted plot
```

```
plot(y=fit_subset$residuals,x=fit_subset$fit,xlab="fitted values",ylab="residuals",main="residual vs fitted")
```



There does not appear to be any strong pattern to the residuals, and they all appear to lie randomly in a horizontal band, so this plot supports the assumption that the errors are independently, identically distributed. The plot does not indicate any outliers.